
Autokeras Text Classification - Final Report

Cesar Lopez	Praveen Venugopal	Vansh Narula
cesarhn95	razorvine	narulavansh3
224005226	727006043	128001615

Abstract

Binary classifiers are used to predict the sentiment behind a give sentence (meaningful sequence of words). One of these famous examples is the IMDB (1) model in which a probabilistic approach is used for determining if a movie review is overall positive or negative. A more recent approach to doing sentiment prediction is by using Neural Networks, specially RNNs. This approach is commonly called Natural Language Processing or NLP. By using a bidirectional LSTM we were able to properly obtain similar sentiment prediction accuracy than previous purely statistic models. However, LSTMs are not trivial, they require the raw input to be embedded, and the size and shape of the model will yield different results and accuracies for a given input. Searching through the best embedding sizes and best model hyperparameters for a model given an input is not a trivial task. Neural Automated search within the Auto-keras framework was used such that hyperparameter search for a given input is trivialized. The end result was an auto-keras models that will look through the best hyperparameters for this LSTM model given an input dataset.

1 Introduction

1.1 Sentiment Analysis

Sentiment Analysis is the area of linguistics that categorizes a statement having a positive or negative connotation. A simple binary classifier, although seen simple and commonly called the Hello World of NLP, sentiment prediction has a wide range of uses ranging from email spam filter to performing automated review de-noising. Using examples from the IMDB dataset (1) This problem might seem trivial from clear examples like:

- [POS] This movie 's AWESOME
- [POS] TRUE ART !!!
- [NEG] Detestable Kubrick social commentary I hold in high disdain .
- [NEG] Stupid , stupid , stupid

This example can be properly categorized by recognizing keywords which have positive or negative connotation. For example words like "Art" and "Awesome" are positive, while "detestable" "stupid" are negative. This model however is easy to fool and incredibly unreliable. This was the premise for Weizenbaum's ELIZA which would suggest the context and the word topics based on keywords. One can see how quickly this example can be turned on its head by the following additions:

- [NEG] This movie 's not AWESOME (*double negative*)
- [NEG] Wannabe Art, pretentious (*sarcasm, ironic tone*)
- [POS] This situation is Detestable, how did Kubrick did not do social commentary before? (*negative terms are abundant*)

- [POS] "Momma says stupid is what stupid does" hahaha (*New term that only belongs to this movie*)

We can see that by quick additions of words, one can flip the polarity of any given statement. As such sentiment analysis is impossible to handcraft. Enter Machine learning. Probabilistic models range from Bag of Words (introduced in the 50s by Zellig Harris) to more complex machine learning protocols that parse through an entire document and quantize parameters between words like the one introduced by Maas et. all (1). For example, the model from Maas et all. scans the documents and extracts *semantic similitatities* then word sentiment then trains a model that maximizes prediction accuracy. In order to do proper testing Maas et all (1) created the IMDB dataset (which we used for our initial testing).

After the famous Khrizhevzy result (2) a revival and the availability of more data, Neural Network based mechanisms (in particular those relying on Neural Networks) have been applied to natural language problems like sentiment analysis. Although initially designed for image processing and categorization, Neural Networks like recursive neural networks based on long short-term memory blocks (typically called LSTM) have gained popularity within text classification solutions. With their benefits, RNN models come with their drawbacks, depending on the wanted accuracy, some models with different shapes and sizes will not be able to achieve a a wanted accuracy.

Recurrent Neural networks as seen on figure ?? are a special type of convolutional neural network that will feed the output of a layer into the input of the next data point. This is helpful because it grants the network the ability of processing data on a element by element basis with the ability of keeping track of the elements that came before it. As a result a RNN has the ability of granting the network non-volatile memory. This is specially helpful for Language models because we can process each sentence on a word by word basis and our network would keep track of the words that came before it.

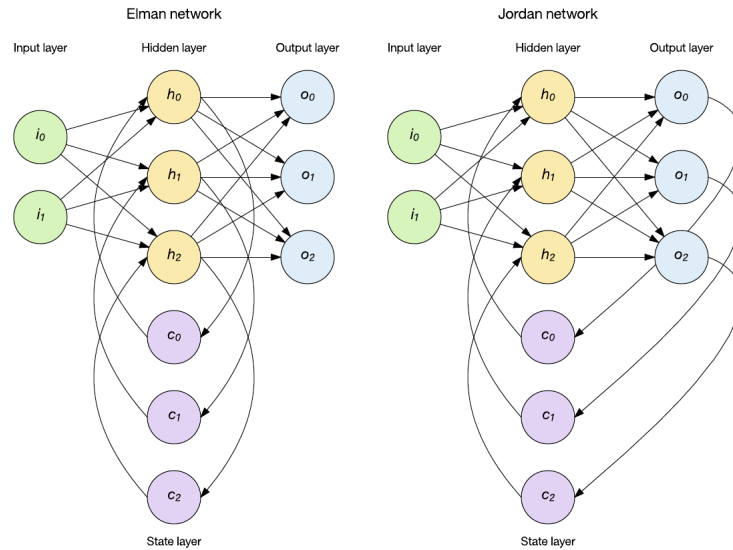


Figure 1: Examples of Recurrent Neural Networks

c

As with other types of neural networks, RNN expect an input on a specific format with size. Given the variation in the length and possible characters in words, it is not possible to feed raw words into an RNN. Word Embeddings such as Words2Vec were initially introduced by Milokov et al. (3) as a way to transform words into a uniform vector of their features, so they can fed into an RNN. As such, word embeddings are a make it or break it part of neural networks.

1.2 Automated Machine Learning

With the new-found popularity of machine learning mechanisms, applications are growing faster than data scientists available for executing them. Automated machine learning (or AutoML) aims

68 at abstracting the intricacies of different machine learning models from the final user. But, not all
69 models will be optimal for all types of data. AutoML approaches this by doing neural automated
70 search. The end goal of NAS is to perform a search inside a range of given tunable hyperparameters
71 to get the model that is optimal for a given data.

72 Given the popularity of the Keras framework in CNN-based machine learning models, Jin et al
73 (4) created Auto-Keras as an AutoML framework based on Keras. Auto-keras aims at further
74 democratizing the power of Machine Learning by being a completely open source framework. This
75 allows the user to create and refine a model from the privacy of their local system. Although other
76 popular closed source AutoML projects like Google's Cloud AutoML have been very popular, they
77 are not compatible with applications in where sensitive data cannot be sent and processed in the
78 cloud. In addition, having the comfort to create an optimal model for your data from your local
79 system allows the model to run and retrain without the availability of an internet connection. The
80 applications are endless.

81 In this work we will aim at creating a generic Binary Sentiment classifier that can identify the correct
82 sentiment of a statement and include it as a supervised class for the Auto-Keras framework. In this
83 way, once included in a framework, our model will search for the best hyperparameters given a generic
84 input and generate the best model for it.

85 2 Related Work

86 2.1 Early Approaches

87 The nuances associated with the term Machine Learning, paints an impression that these problems
88 are quite recent. But the fact would be that such challenges existed for decades and novel approaches
89 have been suggested to tackle them. One of the earliest work on text classification applied a Naive
90 Bayes approach for classification of sentences. Andrew McCallum et al. (5) use a multinomial
91 model to capture word frequencies and calculate probability of occurrence of subsequent words. The
92 paper used Precision/Recall as performance measure. Further advances were done by Kamal et al.
93 in 1999 (6) using Expectation Maximization, and by Simon Tong et al. in 2001 (7) using Support
94 Vector machines. These papers give us a rich insight to the prior work done with respect to Text
95 classification under more pristine settings when Neural Networks were still not in the big screen.

96 The original plan for doing text classification submitted on previous project proposal consisted of
97 applying a *Universal Language Modeling* for text classification. This approach consists on pre-
98 training a model with a large amount of generic data. In this way an initial pre-trained model will
99 contain information about what connotation each word evokes (can vary from binary class to k
100 classes). Once the model is trained for every possible word, the model is fine tuned with data specific
101 to the task we want our model to be tested on.

102 After fine tuning, training is redone by doing *gradual unfreezing*. On this technique our new model is
103 retrained one layer at a time in order to train our model from more general to least general. This is
104 the only approach that the authors mention as adaptable for other models.

105 Although proven to be effective, this model requires access to the original words and cross referencing
106 them with the model that we are going to train (IMDb dataset in our case). This does not align with
107 the mode of operation for auto-keras (4) in where the model has to have little or no information about
108 the data it processes. The ideal auto-keras text classifier will encode the words into different integer
109 IDs in each sentence and our model should be able to draw data from it. As such our Technical plan
110 had to be adapted into these requirements.

111 2.2 Deep Learning approaches

112 Roughly a decade later, Xavier Glorot et al. (?) submitted their work wherein they implemented
113 a classifier using denoising Autoencoders. Post this there were many experimentation done with
114 different variant of Neural networks which included CNN, RNN and RCNN. Recurrent layers showed
115 good results in their field of text classification as proved by Pengfeng et al (8) and Chunting et al (9).
116 This led to a phase where many papers experimented with variants of RNN such as GRUs and LSTM.
117 These papers although showed good accuracy and results, did not introduce any staggering changes

118 to the overall approach. CNNs made a comeback in 2016 in the work of Alexis Conneau et al (10),
119 but soon were overshadowed by recurrent networks.

120 2.3 Universal Language Modelling - ULM

121 While Deep Learning models have achieved state-of-the-art on many NLP tasks,
122 these models are trained from scratch, requiring large datasets, and days to con-
123 verge. Research in NLP focused mostly on transductive transfer (Blitzer et al.,
124 2007)[https://www.cs.jhu.edu/~mdredze/publications/sentiment_acl07.pdf]. For inductive
125 transfer, fine-tuning pre-trained word embeddings (Mikolov et al., 2013), a simple transfer technique
126 that only targets a model's first layer, has had a large impact in practice and is used in most
127 state-of-the-art models. Recent approaches that concatenate embeddings derived from other tasks
128 with the input at layers. This proves to the effectiveness of using pre-trained models. In our work,
129 we combine the proven abilities of LSTM models with pre trained word embeddings to achieve a
130 good accuracy on the given dataset.

131 3 Methodology

132 3.1 FC layers and CNN

133 Our first approach to sentiment prediction was performing word embeddings with our corpus and
134 creating the neural network that precedes it from scratch. The initial plan consisted on implementing
135 a Universal Language Model using LSTM. Although this model is novel and achieves great perfor-
136 mance, its implementation was non-trivial. The first phase pertaining to general domain pre-training
137 of the Language Model requires huge train data set, at least comparable in size with Wikitext-103.
138 Although initially a type of word embedding using TF-IDF/One hot encoding was tried, another
139 large dataset had to be used for initial embedding and it had to share encoding with the targeted data
140 (IMDb in our case) and this was something non feasible with the given data under the time. Given
141 that the course project needed to pass the benchmark set by auto-keras sample train and test data sets
142 which are rather smaller in size, pre-training on those data large sets was deemed non feasible.

143 Parting from UMLFit, we tried to regularize each sentence by turning each word into a one-hot
144 encoded vector and with a vector the size of the maximum amount of words, we fed it into a vanilla
145 fully connected layer, as this was the topic we were learning in class at the moment. Initial testing,
146 even after hyperparameter search showed that the accuracy was not getting higher than 50%, which
147 mean that we might as well be giving random noise to our neural network because it couldn't learn.
148 We had not yet realized that this approach does not translate well because there is no amount of
149 recurrence and the result of processing a word will have no effect on the word processed after it.

150 The next stage we tried was using Convolutional neural networks. We did one hot encoding for the
151 entire sentence. We started this time by using the raw dataset and manually building a vocabulary
152 that mapped an ID to each new word. Then we one-hot encoded a sentence to a 2D array of size $m \times n$
153 where m is the maximum number in the vocabulary and n is the maximum length of a sentences,
154 each n position corresponding to a word will have 1 for the position of the ID its word belonged. Then
155 we connected this to a vanilla CNN. This approach was tried when we were learning CNNs in our
156 class. Although this seemed promising initially, because we kept the encoded words in their order
157 and the words were mapped one to one with our created dictionary, this also did not work, again
158 achieving an accuracy of 50%. The One-hot matrix created was incredibly sparse and any possible
159 data embedded in the form of an ID matching to a word was quickly diluted vanished after doing
160 convolution with neighbors full of zeroes.

161 After our first two tries and knowing that we were not making any progress, and training could not
162 be done in a generic encoding or manually encoding setting, it was evident that the first step of our
163 algorithm had to be finding a usable encoding scheme. This also coincided with the time in class that
164 we started learning about Recurrent Neural Network.


165 3.2 Multi layer Bi-LSTM Model

Post milestone One, our experimentation with Recurrent Neural Network began picking up pace. As
we saw earlier that RNN has a structure to capture time-series data, it still is not efficient in retaining

the information into the long run. Especially for long sentences this drawback hurt performance drastically. Hence we use a specific variant called Long Short Term memory. Standard RNNs suffer from the vanishing gradient problem. LSTMs overcome this by having an extra recurrent state called a cell, c - which can be thought of as the "memory" of the LSTM - and they use multiple gates which control the flow of information into and out of the memory. For more information, go here. We can simply think of the LSTM as a function of x_t , h_t and c_t , instead of just x_t and h_t .

$$(h_t, c_t) = LSTM(x_t, h_t, c_t)$$


Thus, the model using an LSTM looks something like:



The initial cell state, c_0 , like the initial hidden state is initialized to a tensor of all zeros. The sentiment prediction is still, however, only made using the final hidden state, not the final cell state, i.e. $\hat{y} = f(h_T)$

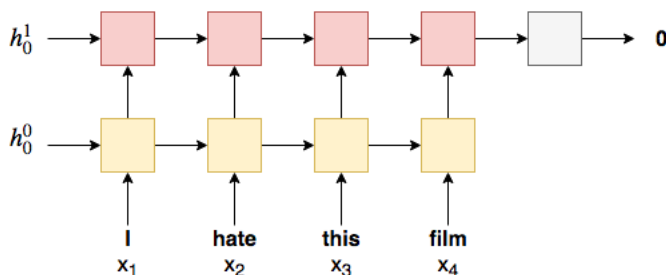
Next we tried with Bi-directional RNN. Herein, in addition to having an RNN processing the words in the sentence from the first to the last (a forward RNN), we have a second RNN processing the words in the sentence from the last to the first (a backward RNN). In PyTorch, the hidden state (and cell state) tensors returned by the forward and backward RNNs are stacked on top of each other. We make our sentiment prediction using the last hidden state from the forward RNN (obtained from final word of the sentence), h_T^{\rightarrow} , and the last hidden state from the backward RNN (obtained from the first word of the sentence), h_T^{\leftarrow} , i.e. $\hat{y} = f(h_T^{\rightarrow}, h_T^{\leftarrow})$

The image below shows a bi-directional RNN, with the forward RNN in orange, the backward RNN in green and the linear layer in silver.



Multi-layer RNNs (also called deep RNNs) are another simple concept. The idea is that we add additional RNNs on top of the initial standard RNN, where each RNN added is another layer. The hidden state output by the first (bottom) RNN at time-step t will be the input to the RNN above it at time step t . The prediction is then made from the final hidden state of the final (highest) layer.

The image below shows a multi-layer unidirectional RNN, where the layer number is given as a superscript. Also note that each layer needs their own initial hidden state, h_0^L .



189 3.3 Pre-trained Word embeddings

190 After successful understanding and implementation of the use cases of LSTM, or RNN in general,
 191 for time series data, we realized that although this helped increasing the accuracy a bit but changing
 192 the architecture alone would not yield the best results. The network alone was not able to learn the
 193 contextual semantics of the sentences which is the key for a good sentiment classification algorithm.
 194 Hence, post this we dedicated our time to fetch good vector representation of our initial raw data.
 195 We came up with an option to use pre-trained word embeddings. Now, instead of having our word
 196 embeddings initialized randomly, they are initialized with some pre-trained vectors. The theory is
 197 that these pre-trained vectors already have words with similar semantic meaning close together in
 198 vector space, e.g. "terrible", "awful", "dreadful" are nearby. This gives our embedding layer a good
 199 initialization as it does not have to learn these relations from scratch.

200 While researching about the same, we came across GloVe : Global Vectors for Word Representations,
 201 which is an unsupervised learning algorithm for obtaining vector representations for words. Training
 202 is performed on aggregated global word-word co-occurrence statistics from a corpus, and the resulting
 203 representations showcase interesting linear substructures of the word vector space.

204 The GloVe model is trained on the non-zero entries of a global word-word co-occurrence matrix,
 205 which tabulates how frequently words co-occur with one another in a given corpus. Populating this
 206 matrix requires a single pass through the entire corpus to collect the statistics. For large corpora,
 207 this pass can be computationally expensive, but it is a one-time up-front cost. Subsequent training
 208 iterations are much faster because the number of non-zero matrix entries is typically much smaller
 209 than the total number of words in the corpus.

210 We used "glove.6B.100d" vectors". 6B indicates these vectors were trained on 6 billion tokens and
 211 100d indicates these vectors are 100-dimensional. Now, After obtaining our word representations, we
 212 pass them as inputs to our Bi-Directional LSTM network. This resulted in a tremendous boost in our
 213 model prediction accuracy. The accuracy boost was solely because the pretrained vectors helped the
 214 model to learn the contextual relationships much better and thus, Now, the model can classify the
 215 sentence using the hidden sentimental relationship between the words in sentences rather than the
 216 words itself.

217 3.4 Hyperparameters

218 Once our LSTM model was properly connected to the datasets given, we proceeded to doing
 219 NAS. For hyper-parameters, we saw initially that changing the maximum vocabulary will reduce
 220 the score at lower numbers while increasing compute time for higher ones. Hence we selected the
 221 maximum vocabulary size as our first tunable hyperparameter. Number of Epochs, hidden dimensions,
 222 dropout and number of layers, are the most common hyperparameters changed in order to increase
 223 performance. Performance was a function of accuracy. Independent of the search algorithm used, the
 224 score, along as a set of hyperparameters used to obtain it was stored. When hyperparameter search was
 225 exhausted or the time limit was surpassed, the fit function selects and prints the best hyperparameters
 226 and stores them internally inside the text classifier class. The range for each hyperparameter are given
 227 in table 1. The colloquially coined "Grad Student Search" was used initially to find these parameter
 228 (which consists on trying different options and then manually picking the best ones.

hyperparameter	range
hidden dimensions	64, 128, 256
batch size	16, 32, 64
number of epochs	2, 3, 5
maximum vocabulary size	20000, 25000
number of hidden layers	2, 3, 4
dropout	0.4, 0.5, 0.6

Table 1: Range for final Greedy search

229 3.5 Grid Search

230 The simplest way of performing hyperparameter optimization has been grid search, or a parameter
231 sweep, which is simply an exhaustive searching through a manually specified subset of the hyperpa-
232 rameter space of a learning algorithm. A grid search algorithm must be guided by some performance
233 metric, in our case prediction accuracy in a held-out validation set. Since the parameter space of a
234 machine learner may include real-valued or unbounded value spaces for certain parameters, manually
235 set bounds and discretization may be necessary before applying grid search.

236 Grid search suffers from the curse of dimensionality, but is often embarrassingly parallel because
237 typically the hyperparameter settings it evaluates are independent of each other.

238 3.6 Random Search

239 This is another simple way of choosing hyperparameters where we randomly generate indexes for
240 all the possible hyperparameters one after another. Then we best the model with the obtained set
241 of hyperparameters and does the same after the training stops for the previous model. We keep a
242 track of the model parameters and the corresponding validation accuracy and thus make sure that
243 the same parameters are not repeated again and again. Thus, training with random combination of
244 hyperparameters at each training cycle until all the possible combinations have been tested or the time
245 limit is exceeded. In this we hope that there is chance that our best model will get trained before the
246 time limit is exceeded and thus, this relies on the probability of selecting the best model out of many
247 possible models. This probability becomes very small as the search space becomes larger. But we can
248 still rely on the fact that the model will somehow reach near to best model parameters in the given
249 time limit. Thus, sometimes random search gives better result than grid search when either search
250 space is very large or the dimensions of the grid is huge because there are many tunable parameters.

251 3.7 Naive Greedy Search

252 To complete the final step of the project, we needed to come up with a smarter search algorithm
253 than grid or random search. So, we came up with a Naive Greedy Search which is a combination
254 of Grad student search and Greedy search where we do not provide the entire space to the greedy
255 search. Rather, we stick to a smartly defined search space based on our understanding of the network
256 and then we proceed with greedy search in that space. Hence, doing this, we are able to reduce the
257 search space significantly. This allows us to traverse through the solutions that are most relevant
258 without going over all the possible hyperparameters like in grid search. For each hyperparameter
259 we traverse through the smartly defined range and select the hyperparameter which gives the best
260 validation accuracy. Once we have the hyperparameter in hand, we fix its value and move onto the
261 next hyperparameter. This way, we perform a naive greedy search for each tunable hyperparameter
262 where the sequence of the hyperparameters is also defined using the Grad student approach where we
263 mimic human thinking process while dealing with hyperparameters and its sequence of selection
264 using greedy search.

265 4 Results

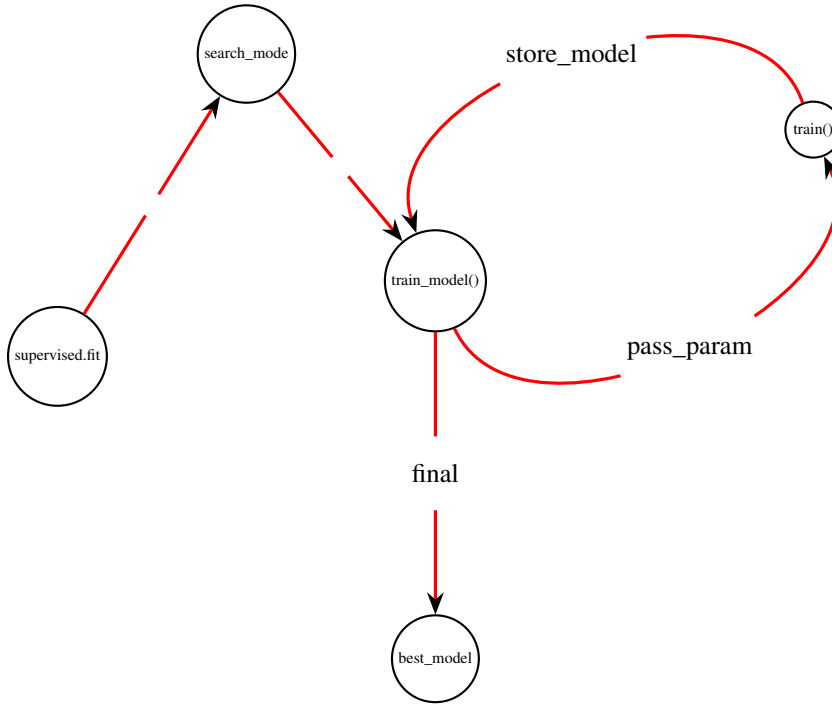
266 4.1 Code Walkthrough

267 Our current implementation uses torchtext framework to load dataset and iterate over it. The code
268 dynamically generates tsv files from the given numpy train data. It is this tsv file that is loaded as
269 torchtext dataset, and iterated over. Two global torchtext data fields are created one to hold the input
270 sentence and the other to hold output sentiment of the train data. We build vocabulary for the data
271 fields only using the train data, during model training.

272 Word-Embeddings are another crucial aspect of this section. We use Glove.6B package with a size of
273 100-dimensions. This converts the individual word in the sentence to a vector of 100-D. Consequently
274 our embedding_layer has a size of 100, throughout the project.

275 Finally, each search algorithm is either exhaustive (will not repeat a search) and has a soft stop for a
276 time limit passed. This means that once the time runs out, fit will finish running the model that it is
277 training before searching for the best one. Once the search is done, fit will select the best model from

the entire search history, even when doing grid search. This ensures that even if a bad decision is made by Greedy search, we can always go back to our best previous state.



4.2 Initial Results with final model

In order to do initial testing, we limited our runs to do the small 96k + 4k raw dataset. After verifying that we could train and keep different models, we selected a 332 configuration. Given that grid search and random search are lengthy to exhaust, we limited our search space to 3 values of hidden dimensions, 3 values of batch sizes and 2 values of number of layers. With this configuration, both grid and random will calculate $6 \times 6 \times 3$, models with a total of 18, and greedy will calculate $3 + 3 + 2 = 8$. We can see that Greedy converges faster, although it is not exhaustive. The progression of these algorithms on a model by model can be seen in tables 5, 5 and 5 (present in the appendix). We can see on table 5 that the random search is exhaustive (no repeated values) and each new set of values is selected in a random order. for grid in table 5 we can see that we explore every possible choice of each value. Finally, for greedy search 5 we can see that we selected the best hyperparameter for hidden dimensions, with this set, we went on and searched for the best batch size and finally the same for number of hidden dimensions.

4.3 Final Results with final model

Once we were comfortable that Greedy search was as efficient as the other two but with an exponentially smaller run time (due to its smart reduced search space) a 3-3-2-3-3 search was done. For all parameters except max vocab size, we selected 3 possible values. Having set our final search space, we proceeded to get the best hyperparameters through grid search for both the small raw dataset (96k +4k) and the raw large dataset (480k +). The results of the run for the small dataset are shown in the appendix as table 5. The final results for the large dataset are shown in table 4.3. In here we can see that the greedy in runs 1,2,3 converges that best accuracy is obtained at 3 layers, then in runs 4,5,6 it learns that the best number of hidden dimensions is 256 and the same thing happens for the other hyperparameters.

Run	Accuracy	hidden_dim	batch	epoch	vocab	layers	dropout	time
1	83.65	128	16	2	20000	2	0.4	418.58
2	83.70	128	16	2	20000	3	0.4	593.71
3	83.67	128	16	2	20000	4	0.4	847.53
4	83.43	64	16	2	20000	3	0.4	515.43
5	83.54	128	16	2	20000	3	0.4	594.29
6	83.68	256	16	2	20000	3	0.4	929.74
8	83.63	256	16	2	25000	3	0.4	943.97
9	83.58	256	16	2	20000	3	0.4	926.80
10	83.48	256	32	2	20000	3	0.4	684.94
11	82.81	256	64	2	20000	3	0.4	546.93
12	83.76	256	16	2	20000	3	0.4	930.28
13	84.44	256	16	3	20000	3	0.4	1394.24
14	84.80	256	16	5	20000	3	0.4	2318.21
15	84.88	256	16	5	20000	3	0.4	2314.91
16	84.18	256	16	5	20000	3	0.5	2306.06
17	83.46	256	16	5	20000	3	0.6	2301.99

Table 2: Results for greedy search 5 tunable parameters, 3|3|2|3|3| options for 480k dataset

5 Conclusion

After days of extensive learning and testing we have come up with a final Multi-layer Bi directional LSTM model that performs Naive Greedy Search for hyper parameter optimization for sentiment classification. This uses pretrained GloVe embedding to learn the contextual relationship between the words in a sentence and then classifies the sentence based on this relation. This model makes use Auto-Keras Supervised class and hence can be included inside the original Auto-Keras code with potential improvements in search strategy and word embedding, which at last is the room for future improvement that is always there in any solution. The final accuracy achieved after the greedy search on the final test data set is 85.4% and the chosen hyperparameters are 256 hidden dimensions, 16 batch size, 5 epochs, vocab size 20k, 3 layers and 0.4 dropout. The parameters and accuracy are not fixed and can be improved with larger search space but that will result in increased search time. this, based on our understanding and skill set, is the most efficient way to automate text classification task in the given time frame.

Work Distribution and Acknowledgement

Each and every member performed and contributed equally to all the tasks. The work was divided but was never completed by anyone alone and hence, the team actually showed great amount of collaborative efforts which towards the end equipped each one with a broader understanding of the task in hand. The whole project was looked upon as an opportunity to grow and learn the basics of NLP and the use cases of Neural Networks in solving such tasks. We are glad the outcome was as expected.

We would like to express our gratitude to our mentor Haifeng Jin and professor Shuiwang Ji for always being there and guiding us throughout this phase of learning. We thank you for providing us with the task and considering us capable enough to complete it in the given time frame. We hope we have done justice to your expectations.

References

- [1] A. L. Maas, R. E. Daly, P. T. Pham, D. Huang, A. Y. Ng, and C. Potts, "Learning word vectors for sentiment analysis," in *Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics: Human Language Technologies*, (Portland, Oregon, USA), pp. 142–150, Association for Computational Linguistics, June 2011.
- [2] A. Krizhevsky, I. Sutskever, and G. E. Hinton, "Imagenet classification with deep convolutional neural networks," in *Proceedings of the 25th International Conference on Neural Information Processing Systems - Volume 1*, NIPS'12, (USA), pp. 1097–1105, Curran Associates Inc., 2012.

- 336 [3] T. Mikolov, I. Sutskever, K. Chen, G. Corrado, and J. Dean, "Distributed representations of words and
337 phrases and their compositionality," *CoRR*, vol. abs/1310.4546, 2013.
- 338 [4] H. Jin, Q. Song, and X. Hu, "Auto-keras: Efficient neural architecture search with network morphism,"
339 2018.
- 340 [5] A. McCallum, K. Nigam, *et al.*, "A comparison of event models for naive bayes text classification," in
341 *AAAI-98 workshop on learning for text categorization*, vol. 752, pp. 41–48, Citeseer, 1998.
- 342 [6] K. Nigam, A. K. Mccallum, S. Thrun, and T. Mitchell, "Text classification from labeled and unlabeled
343 documents using em," *Machine Learning*, vol. 39, pp. 103–134, May 2000.
- 344 [7] S. Tong and D. Koller, "Support vector machine active learning with applications to text classification," *J.*
345 *Mach. Learn. Res.*, vol. 2, pp. 45–66, Mar. 2002.
- 346 [8] P. Liu, X. Qiu, and X. Huang, "Recurrent neural network for text classification with multi-task learning,"
347 *CoRR*, vol. abs/1605.05101, 2016.
- 348 [9] C. Zhou, C. Sun, Z. Liu, and F. C. M. Lau, "A C-LSTM neural network for text classification," *CoRR*,
349 vol. abs/1511.08630, 2015.
- 350 [10] A. Conneau, H. Schwenk, L. Barrault, and Y. LeCun, "Very deep convolutional networks for natural
351 language processing," *CoRR*, vol. abs/1606.01781, 2016.

352 Appendix

Accuracy	hidden_dim	batch	epoch	vocab	layers	dropout	time
76.00	64	4	1	25000	2	0.5	140.29
75.95	128	4	1	25000	2	0.5	142.93
76.05	256	4	1	25000	2	0.5	174.57
76.32	256	4	1	25000	2	0.5	171.85
74.94	256	8	1	25000	2	0.5	98.93
75.12	256	12	1	25000	2	0.5	74.59
76.43	256	4	1	25000	2	0.5	174.60
76.10	256	4	1	25000	3	0.5	248.21
76.23	256	4	1	25000	2	0.5	173.46
75.28	256	4	1	25000	3	0.5	246.35
75.89	256	8	1	25000	2	0.5	99.94
75.20	256	8	1	25000	3	0.5	146.14
74.82	256	12	1	25000	2	0.5	74.92
75.13	256	12	1	25000	3	0.5	112.04

Table 3: Results for grid search 3 tunable parameters, 31312 options

353 **Random Search**

Accuracy	hidden_dim	batch	epoch	vocab	layers	dropout	time
79.41	128	16	2	20000	2	0.4	86.73
79.08	128	16	2	20000	3	0.4	119.62
79.35	128	16	2	20000	4	0.4	169.30
79.30	64	16	2	20000	2	0.4	78.50
79.77	128	16	2	20000	2	0.4	85.07
79.68	256	16	2	20000	2	0.4	123.00
79.83	128	16	2	20000	2	0.4	85.41
79.08	128	16	2	25000	2	0.4	88.43
79.37	128	16	2	20000	2	0.4	82.77
78.30	128	32	2	20000	2	0.4	54.00
76.92	128	64	2	20000	2	0.4	39.30
79.27	128	16	2	20000	2	0.4	85.55
80.66	128	16	3	20000	2	0.4	126.40
81.76	128	16	5	20000	2	0.4	208.95
81.82	128	16	5	20000	2	0.4	208.94
81.08	128	16	5	20000	2	0.5	208.97
80.21	128	16	5	20000	2	0.6	209.16

Table 4: Results for random search 3 tunable parameters, 31312 options

354 **Greedy Search**

Accuracy	hidden_dim	batch	epoch	vocab	layers	dropout	time
76.00	64	4	1	25000	2	0.5	140.29
75.95	128	4	1	25000	2	0.5	142.93
76.05	256	4	1	25000	2	0.5	174.57
76.32	256	4	1	25000	2	0.5	171.85
74.94	256	8	1	25000	2	0.5	98.93
75.12	256	12	1	25000	2	0.5	74.59
76.43	256	4	1	25000	2	0.5	174.60
76.09	256	4	1	25000	3	0.5	248.21

Table 5: Results for greedy search 3 tunable parameters, 31312 options

Accuracy	hidden_dim	batch	epoch	vocab	layers	dropout	time
76.37	128	4	1	25000	2	0.5	143.75
73.74	64	8	1	25000	2	0.5	72.94
75.57	128	12	1	25000	2	0.5	53.48
74.53	64	12	1	25000	3	0.5	66.62
75.59	64	12	1	25000	2	0.5	50.92
74.82	128	12	1	25000	3	0.5	73.26
75.72	64	8	1	25000	3	0.5	93.97
76.07	64	4	1	25000	2	0.5	138.90
75.98	128	4	1	25000	3	0.5	184.14
75.20	256	8	1	25000	3	0.5	146.13
75.38	128	8	1	25000	3	0.5	100.64
76.23	256	4	1	25000	2	0.5	173.49
75.27	256	4	1	25000	3	0.5	246.34
75.88	256	8	1	25000	2	0.5	99.94
74.81	256	12	1	25000	2	0.5	74.92
76.08	64	4	1	25000	3	0.5	175.36
75.40	128	8	1	25000	2	0.5	75.98
75.12	256	12	1	25000	3	0.5	112.03

Table 6: Results for greedy search 6 tunable parameters, 3|3|2|3|3|3 options for 96k+4k dataset